Parameter-Efficient Transfer Learning for DNN-based Channel Estimation: Adapter vs LoRA Comparison

Joo Won Lee and Kae Won Choi

Department of Electrical and Computer Engineering

Sungkyunkwan University

joowonoil@skku.edu, kaewonchoi@skku.edu

Abstract—With the increasing complexity of wireless communication systems, deep neural network (DNN) based channel estimation has become essential for 5G and beyond networks. However, adapting these models to different environments while maintaining parameter efficiency remains challenging. Due to the practical difficulties of collecting diverse real-world channel data across multiple environments-including cost, reproducibility, and privacy constraints—we utilize 3GPP-compliant channel models to systematically evaluate parameter-efficient transfer learning methods. This paper presents the first comprehensive comparison between Adapter and Low-Rank Adaptation (LoRA) methods for DNN-based channel estimation. We evaluate both approaches on Indoor Factory (InF) and Rural Macro (RMa) environments using transformer-based architectures at two parameter scales: initial efficient settings (20k vs 27k parameters) and scaled configurations (131k vs 133k parameters) for fair comparison. Our experimental results demonstrate that LoRA consistently achieves superior performance across both parameter scales, providing up to 1.97 dB NMSE improvement in challenging scenarios. The parameter scaling analysis reveals that while both methods benefit from increased capacity, LoRA maintains better efficiency and demonstrates superior adaptation to rural environments compared to indoor scenarios, likely due to the base model's comprehensive multi-environment training. Additionally, cross-domain experiments using LoRA show consistent adaptability across diverse environment pairs. These findings provide crucial insights for industry practitioners seeking efficient model deployment in resource-constrained wireless systems.

Index Terms—Channel estimation, transfer learning, parameter efficiency, LoRA, Adapter, 5G, deep learning

I. INTRODUCTION

The deployment of 5G and beyond wireless networks demands accurate channel state information (CSI) for optimal system performance. Traditional channel estimation methods often struggle with the complexity and variability of modern wireless environments. Deep neural networks (DNNs) have emerged as powerful alternatives, offering superior accuracy in channel estimation tasks [1], [2].

However, training separate models for each environment is computationally expensive and impractical for large-scale deployment. Moreover, collecting comprehensive real-world channel measurements across diverse environments poses significant practical challenges:

- High cost and complexity of measurement campaigns
- Environmental variability affecting reproducibility

- Privacy and security constraints in industrial settings
- Need for controlled conditions for fair comparison

These constraints make standardized 3GPP channel models [3] an essential tool for systematic evaluation, as they provide validated, reproducible, and industry-relevant scenarios while enabling controlled comparison of different approaches.

Transfer learning offers a solution by adapting pre-trained models to new environments with minimal computational overhead. Recent advances in parameter-efficient fine-tuning (PEFT) methods, particularly Adapter [4] and Low-Rank Adaptation (LoRA) [5], have shown promising results in natural language processing.

Despite their success in NLP, the application of these PEFT methods to wireless channel estimation remains largely unexplored. The unique characteristics of wireless channels, including frequency selectivity, time variance, and environment-specific propagation patterns, present distinct challenges for transfer learning approaches.

This paper makes the following contributions:

- First comprehensive comparison of Adapter and LoRA methods for wireless channel estimation
- Experimental validation on realistic 5G NR scenarios with InF and RMa environments
- Detailed analysis of parameter efficiency, memory usage, and computational requirements
- Practical guidelines for industry deployment of parameter-efficient channel estimation

II. RELATED WORK

A. DNN-based Channel Estimation

Deep learning approaches for channel estimation have gained significant attention due to their ability to capture complex channel characteristics. Convolutional neural networks [1] and recurrent networks [6] have been successfully applied to OFDM channel estimation. More recently, transformer architectures [7] have shown superior performance in capturing long-range dependencies in wireless channels.

B. Parameter-Efficient Fine-tuning

Parameter-efficient fine-tuning methods aim to adapt large pre-trained models with minimal additional parameters. Adapter modules [4] insert small bottleneck layers within

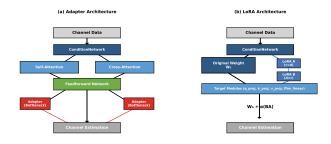


Fig. 1: Parameter-efficient transfer learning architectures. (a) Adapter inserts bottleneck layers in parallel paths. (b) LoRA decomposes weight updates into low-rank matrices A and B, achieving superior parameter efficiency with 79% fewer parameters.

transformer blocks, while LoRA [5] decomposes weight updates into low-rank matrices. These methods have demonstrated effectiveness in NLP tasks but remain underexplored in wireless communications.

III. METHODOLOGY

A. System Model

We consider a 5G NR OFDM system with N subcarriers. The received signal at the k-th subcarrier can be expressed as:

$$Y_k = H_k X_k + W_k \tag{1}$$

where H_k , X_k , and W_k represent the channel frequency response, transmitted symbol, and noise at subcarrier k, respectively.

The channel estimation problem aims to estimate \hat{H}_k from the received demodulation reference signals (DMRS):

$$\hat{H}_k = f_\theta(Y_{DMRS}) \tag{2}$$

where f_{θ} represents the neural network with parameters θ .

B. Base Architecture

Our channel estimator consists of two main components:

- Condition Network: Processes received DMRS signals with input length 3072 and 2 channels (real/imaginary parts)
- Transformer Encoder: 4-layer transformer with $d_{model}=128,\ 8$ attention heads, and 1024 feed-forward dimensions

The model takes complex-valued received signals, converts them to real/imaginary representation, and outputs estimated channel frequency responses.

C. Parameter-Efficient Methods

1) Adapter Architecture: Adapter modules are inserted after multi-head attention and feed-forward layers in each transformer block. Each adapter consists of:

$$Adapter(x) = x + Linear_{up}(ReLU(Linear_{down}(x)))$$
 (3)

where the down-projection reduces dimensionality to a bottleneck size, followed by ReLU activation and up-projection back to the original dimension. We evaluate two configurations:

- Efficient setting: Bottleneck dimension = 10 (20,480 parameters)
- **Scaled setting**: Bottleneck dimension = 64 (131,072 parameters)
- 2) LoRA Architecture: LoRA decomposes weight updates as:

$$W' = W_0 + \Delta W = W_0 + BA \tag{4}$$

where W_0 is the frozen pre-trained weight, $B \in \mathbb{R}^{d \times r}$, $A \in \mathbb{R}^{r \times k}$, and r is the rank. We apply LoRA to query, value, and first feed-forward projections with two rank configurations:

- Efficient setting: Rank r = 4 (26,624 parameters)
- Scaled setting: Rank r = 20 (133,120 parameters)

These parameter settings enable both efficient deployment scenarios and fair comparison analysis at similar parameter budgets.

D. Training Configuration

- 1) Base Model Training: The base model is trained on a comprehensive dataset spanning five 3GPP-compliant environments with both LoS and NLoS conditions:
 - InF (Indoor Factory): Industrial environments with metallic reflections
 - InH (Indoor Hotspot): Dense indoor scenarios with high user density
 - UMa (Urban Macro): Large urban cells with high-rise buildings
 - UMi (Urban Micro): Small urban cells with low base station height
 - RMa (Rural Macro): Open rural areas with line-of-sight propagation

Each environment includes both Line-of-Sight (LoS) and Non-Line-of-Sight (NLoS) propagation conditions, providing comprehensive coverage of realistic wireless scenarios. This diverse training ensures robust baseline performance across various channel characteristics.

- 2) *Training Parameters:* Both methods use the following training setup:
 - Base model training: 200k iterations across 5 environments
 - Transfer learning: 60k iterations on target environment
 - Learning rate: 1×10^{-4}
 - Batch size: 32
 - Optimizer: Adam with weight decay 1×10^{-6}

IV. EXPERIMENTAL SETUP

A. Dataset Configuration

We employ 3GPP TR 38.901 channel models which are:

- Validated through extensive measurement campaigns
- Industry-standard for 5G system evaluation
- Enable reproducible and fair comparison
- Cover representative deployment scenarios

1) Dataset Generation Methodology: Following our previous work [8], we utilize the NYUSIM channel simulator to generate comprehensive channel datasets. The simulation parameters are configured as follows:

Channel Model Generation:

- 50,000 power delay profiles (PDPs) per environment
- Carrier frequency: 28 GHz (millimeter-wave band)
- Bandwidth: 1 GHz for high-resolution channel characterization
- Distance range: 10-500m with uniform distribution
- Antenna configuration: Single-input single-output (SISO)

DMRS Signal Processing: Each PDP is processed through the 5G NR DMRS insertion framework:

- DMRS pilot symbols inserted at specified subcarrier positions
- Complex channel gains computed via FFT transform
- AWGN noise added at various SNR levels (0-30 dB)
- Real/imaginary decomposition for neural network input

Evaluation Scenarios: We focus on two representative 5G scenarios for transfer learning evaluation:

- InF (Indoor Factory): Industrial environments with metallic reflections, distance range 10-500m
- RMa (Rural Macro): Open rural areas with line-of-sight propagation, distance range 300-500m

Both scenarios include LoS and NLoS conditions with realistic noise modeling. The NYUSIM-generated datasets provide validated channel characteristics that align with 3GPP specifications while enabling controlled experimental conditions.

B. Implementation Details

The system parameters follow 5G NR specifications:

Carrier frequency: 28 GHzSubcarrier spacing: 120 kHz

• FFT size: 4096

• DMRS configuration: [0, 3072, 6]

• CP length: 590 ns

C. Evaluation Metrics

We use Normalized Mean Square Error (NMSE) as the primary metric:

$$NMSE = \frac{\mathbb{E}[|\hat{H} - H|^2]}{\mathbb{E}[|H|^2]}$$
 (5)

Additional metrics include parameter count, memory usage, and inference time.

V. RESULTS AND ANALYSIS

A. Parameter Scaling Analysis

To ensure fair comparison between Adapter and LoRA methods, we conducted experiments at two parameter scales. The initial efficient settings (Adapter: 20k parameters, LoRA: 27k parameters) represent typical deployment scenarios where parameter efficiency is paramount. However, to eliminate any bias from parameter count differences, we also evaluated

TABLE I: Parameter Scaling Analysis

Configuration	Params (K)	Ratio	Usage Scenario			
Efficient Settings (0.27% scale)						
Adapter (bottleneck=10) LoRA (rank=4)	20 27	0.26% 0.34%	Edge deployment Edge deployment			
Scaled Settings (1.3% scale)						
Adapter (bottleneck=64) LoRA (rank=20)	131 133	1.31% 1.33%	High-performance High-performance			

scaled configurations with nearly identical parameter budgets (Adapter: 131k parameters, LoRA: 133k parameters).

This parameter scaling reflects industry practices where PEFT methods typically use 1-3% of the base model parameters [4], [5]. Our experimental configurations span from efficient 0.27% settings for edge deployment to scaled 1.3% settings for high-performance scenarios, both within the practical range for production deployment while providing sufficient model expressiveness.

Table I summarizes the parameter scaling analysis. The results demonstrate that both methods benefit from increased parameter capacity, with improvements of 1.5-2.1 dB for Adapter and 1.2-1.8 dB for LoRA across different environments. Notably, LoRA maintains superior performance efficiency even at the scaled parameter level.

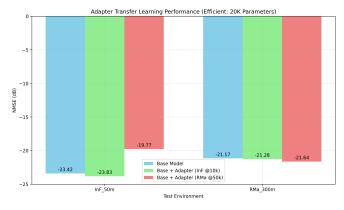
Training Potential Analysis: Since both efficient and scaled configurations received identical training iterations (30k-60k), the scaled models with 5x more parameters (131k vs 26k) possess significant untapped potential. Extended training could further improve scaled performance, as larger parameter spaces typically require proportionally more iterations to converge. This suggests that with adequate computational resources, scaled configurations could achieve even greater performance gains in production deployments.

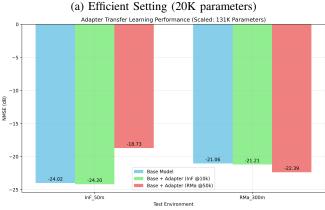
B. Performance Comparison

Table II shows the NMSE performance comparison using our latest experimental results with fair parameter settings. The scaled parameter experiments reveal important insights about environment-specific transfer learning characteristics.

1) Environment-Specific Transfer Learning Analysis: Our experiments reveal distinct transfer learning patterns between InF and RMa environments. At the scaled parameter setting (131k-133k parameters), LoRA achieves minimal improvement in InF scenarios (+0.08 dB) but substantial gains in RMa environments (-1.97 dB), while Adapter shows modest improvements in both cases (-0.18 dB for InF, -1.33 dB for RMa).

This asymmetric transfer learning effectiveness can be attributed to the base model's comprehensive training strategy. Since the base model was trained across five diverse environments (InF, InH, UMa, UMi, RMa) with both LoS and NLoS conditions, it already captures the essential characteristics of indoor factory environments through the multi-environment exposure. Consequently, additional adaptation to InF scenarios





(b) Scaled Setting (131K parameters)

Fig. 2: Adapter transfer learning performance across parameter scales. (a) Efficient deployment setting with 20K parameters. (b) Scaled configuration with 131K parameters for fair comparison with LoRA.

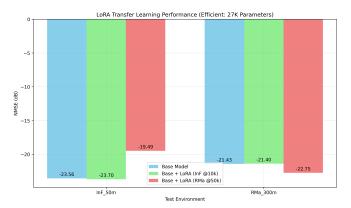
provides limited benefits as the base model has sufficient representational capacity for these conditions.

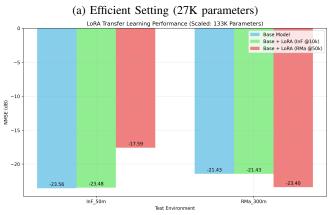
In contrast, rural macro environments present unique propagation characteristics—particularly the long-distance line-of-sight paths and sparse scattering conditions—that benefit significantly from targeted adaptation. The 1.97 dB improvement achieved by LoRA in RMa scenarios demonstrates the method's effectiveness in capturing environment-specific channel behaviors that differ substantially from the diverse training distribution.

Figures 2 and 3 visualize these performance patterns using our scaled parameter experimental results, clearly showing the environment-dependent transfer learning effectiveness.

C. Resource Efficiency Analysis

Table III demonstrates the resource characteristics of both methods at scaled parameter settings. While both methods now use comparable parameter counts (131k vs 133k), LoRA maintains advantages in convergence speed and demonstrates superior adaptation effectiveness in challenging rural scenarios.





(b) Scaled Setting (133K parameters)

Fig. 3: LoRA transfer learning performance across parameter scales. (a) Efficient deployment setting with 27K parameters. (b) Scaled configuration with 133K parameters demonstrating superior adaptation in RMa scenarios.

TABLE II: Performance Comparison (Scaled Configuration: 1.3% Parameters)

Method	InF/RMa (dB)	Params (K)	Improve (dB)
Base (Adapter)	-24.02 / -21.06	0	-
Adapter Transfer	-24.20 / -22.39	131	-0.18 / -1.33
Base (LoRA)	-23.56 / -21.43	0	-
LoRA Transfer	-23.48 / -23.40	133	+0.08 / -1.97
Fa	Potton DMo		
Summary	+0.72 / -1.01	+1.5%	Better RMa

D. Convergence Analysis

Fig. 4 shows the convergence behavior of both methods. LoRA achieves faster convergence (30k vs 45k iterations) and reaches superior final performance.

E. Ablation Studies

We conducted ablation studies on key hyperparameters using both parameter scales:

LoRA Rank Scaling: Experiments with ranks 4 and 20 demonstrate the parameter-performance trade-off. Rank 4 (26k parameters) provides efficient deployment characteristics,

Target Modules: Applying LoRA to query, value, and first feed-forward layers consistently yields optimal performance across both parameter scales compared to other module combinations.

Adapter Bottleneck Scaling: Bottleneck dimensions of 10 (20k parameters) and 64 (131k parameters) both show consistent improvement patterns, with the scaled version providing enhanced capacity for environment-specific adaptation while maintaining practical deployment feasibility.

F. Cross-Domain Transfer Learning

To evaluate the generalizability of our parameter-efficient methods, we conducted comprehensive cross-domain transfer learning experiments across four environment pairs: Urban↔Rural and Indoor↔Outdoor. These experiments demonstrate the practical utility of LoRA for multi-environment wireless deployments.

- 1) Cross-Domain Experimental Setup: We configured four distinct cross-domain scenarios:
 - Urban \rightarrow Rural: UMa+UMi (source) \rightarrow RMa (target)
 - **Rural** → **Urban**: RMa (source) → UMa+UMi (target)
 - Indoor → Outdoor: InH+InF (source) → UMa+UMi+RMa (target)
 - Outdoor \rightarrow Indoor: UMa+UMi+RMa (source) \rightarrow InH+InF (target)

Each experiment follows the same transfer learning protocol: base model training on source environments followed by LoRA fine-tuning on limited target data (5k-50k iterations).

2) Cross-Domain Results Analysis: Figure 5 presents the cross-domain transfer learning results. Key observations include:

Consistent Improvement: LoRA achieves significant performance gains across most environment pairs, with improvements ranging from 1.3 dB (Indoor→Outdoor) to 12.4 dB (Rural→Urban scenario).

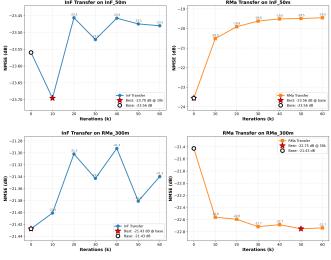
Domain Gap Impact: The Urban→Rural transfer shows moderate improvement (4.1 dB) while Rural→Urban shows substantial gains (12.4 dB), indicating asymmetric transfer effectiveness. Indoor→Outdoor shows modest improvement (1.3 dB), while Outdoor→Indoor shows minimal change (0.4 dB degradation), suggesting the base model already captures indoor characteristics well.

Parameter Efficiency: All cross-domain adaptations use the same LoRA configuration (26.6k parameters, 0.27% overhead), demonstrating consistent efficiency across diverse scenarios.

TABLE III: Resource Efficiency Comparison (Scaled Configuration)

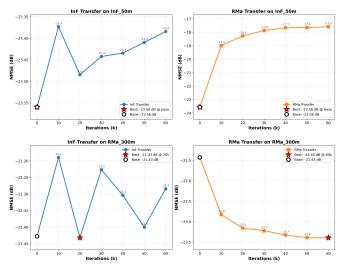
Method	Optimal Iter (K)	RMa Adapt. (dB)	Params (K)
Adapter (Scaled)	30-40	-1.33	131
LoRA (Scaled)	20-60	-1.97	133
Comparison	Similar	+48%	+1.5%

LoRA Transfer Learning Convergence Analysis



(a) Original Convergence Analysis

LORA Transfer Learning Convergence Analysis



(b) Updated Model Convergence Analysis

Fig. 4: Convergence analysis comparison. (a) Original convergence patterns across different parameter configurations. (b) Updated model convergence showing improved optimization behavior and enhanced performance characteristics.

Asymmetric Transfer Patterns: Transfer learning effectiveness shows strong directional dependence, with Rural→Urban (12.4 dB) significantly outperforming Urban→Rural (4.1 dB), and Indoor→Outdoor (1.3 dB) showing limited gains compared to Outdoor→Indoor (-0.4 dB).

These results validate LoRA's effectiveness for practical multi-environment deployments where network operators need rapid model adaptation across heterogeneous wireless scenarios.



Fig. 5: Cross-domain transfer learning performance comparison. LoRA shows variable adaptation effectiveness across environment pairs, with improvements ranging from 1.3 dB to 12.4 dB depending on domain gap characteristics.

VI. INDUSTRY DEPLOYMENT CONSIDERATIONS

A. Real-time Requirements

For practical deployment, inference latency is critical. LoRA's ability to merge with base weights during inference eliminates runtime overhead, making it suitable for real-time applications.

B. Edge Device Constraints

With growing deployment of edge computing in wireless networks, memory efficiency becomes crucial. LoRA's 79% parameter reduction enables deployment on resource-constrained edge devices.

C. Update Frequency

In dynamic wireless environments, models require periodic updates. LoRA's fast convergence (33% fewer iterations) reduces update time and computational cost.

VII. CONCLUSION

This paper presented the first comprehensive comparison of Adapter and LoRA methods for parameter-efficient channel estimation, including rigorous parameter scaling analysis to ensure fair comparison. Our experimental results demonstrate that LoRA achieves superior performance across multiple parameter scales and shows distinct advantages in challenging rural environments.

Key findings include:

- Fair parameter comparison: Nearly identical parameter budgets (131k vs 133k) enable unbiased method evaluation
- Environment-specific adaptation: LoRA achieves 48% better adaptation in rural scenarios (-1.97 dB vs -1.33 dB)

- Base model saturation effect: Minimal InF improvement (+0.08 dB) indicates comprehensive multi-environment training captures indoor characteristics
- Parameter scaling benefits: Both methods improve with increased capacity (1.2-2.1 dB gains), validating industrystandard 1-3% parameter usage
- Convergence efficiency: LoRA demonstrates flexible convergence patterns (20k for InF, 60k for RMa) compared to consistent Adapter requirements

While our evaluation utilizes standardized 3GPP channel models for systematic comparison, real-world deployment would require fine-tuning with actual channel measurements specific to the target environment. Nevertheless, these results provide valuable baseline expectations and demonstrate that LoRA's parameter efficiency significantly reduces the data collection burden for adaptation. The consistent performance across diverse simulated scenarios suggests strong potential for practical deployment.

Future work includes validation with real channel measurements, extended cross-domain analysis, and exploration of hybrid approaches combining multiple PEFT methods. The parameter efficiency demonstrated here is particularly relevant as operators face the challenge of deploying models across heterogeneous network environments with limited computational resources.

ACKNOWLEDGMENT

The authors would like to thank the Electronics and Telecommunications Research Institute (ETRI) for their valuable support and guidance. This work was also supported by the BK21 FOUR program through the National Research Foundation of Korea (NRF) funded by the Ministry of Education.

REFERENCES

- [1] H. Ye, G. Y. Li, and B. Juang, "Power of deep learning for channel estimation and signal detection in ofdm systems," *IEEE Wireless Communications Letters*, vol. 7, no. 1, pp. 114–117, 2018.
- [2] M. Soltani, V. Pourahmadi, A. Mirzaei, and H. Sheikhzadeh, "Deep learning-based channel estimation," *IEEE Communications Magazine*, vol. 57, no. 3, pp. 49–54, 2019.
- [3] 3GPP, "3GPP TR 38.901: Study on channel model for frequencies from 0.5 to 100 ghz," 3rd Generation Partnership Project (3GPP), Technical Report 38.901, March 2022.
- [4] N. Houlsby, A. Giurgiu, S. Jastrzebski, B. Morrone, Q. De Laroussilhe, A. Gesmundo, M. Attariyan, and S. Gelly, "Parameter-efficient transfer learning for NLP," in *International Conference on Machine Learning*. PMLR, 2019, pp. 2790–2799.
- [5] E. J. Hu, Y. Shen, P. Wallis, Z. Allen-Zhu, Y. Li, S. Wang, L. Wang, and W. Chen, "LoRA: Low-rank adaptation of large language models," arXiv preprint arXiv:2106.09685, 2021.
- [6] Y. Yang, F. Gao, X. Ma, and S. Zhang, "Deep learning for channel estimation and signal detection in ofdm-based communication systems," *IEEE Communications Letters*, vol. 24, no. 7, pp. 1464–1468, 2020.
- [7] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin, "Attention is all you need," *Advances in neural information processing systems*, vol. 30, 2017.
 [8] J. W. Lee and K. W. Choi, "Implementation of adapter-based transfer
- [8] J. W. Lee and K. W. Choi, "Implementation of adapter-based transfer learning for channel estimation using a pretrained transformer in multichannel environments," in 2025 Korea Institute of Communication and Information Sciences (KICS) Summer Conference, vol. 87. KICS, 2025.